

Synchronization Choices Overview

► Typical uses/techniques, and **examples** we'll consider:

	Synchronized externally (by the caller)	Synchronized internally	Both
Using locks	Default, esp. when transaction involves multiple objects <i>Producer/Consumer: locks</i>	"Synchronized" objects (Generally problematic)	<i>(e.g., copy-on-write strings)</i>
Using atomics ("lock-free")	Atomic handoffs <i>Producer/Consumer: atomic mail slots</i> <i>Versioned state</i>	"Lock-free" data structures <i>slist</i>	
Both	Different tools at different times <i>Double-Checked Locking (DCL)</i> <i>Producer/Consumer: locks + atomics</i>	<i>config_map</i>	

►

Why Lock-Free Code?

- ▶ **Concurrency and scalability.**

- ▶ Eliminate/reduce blocking/waiting in algorithms and data structures.

- ▶ Avoid the troubles with (b)locking:

```
{  
    lock_guard<mutex> lock1{ mutTable1 };  
    lock_guard<mutex> lock2{ mutTable2 };  
    table1.erase( x );  
    table2.insert( x );  
} // release mutTable2 and mutTable1
```


- ▶ **Races:** Forgot to lock, or locked the wrong thing.
- ▶ **Deadlock:** Locked in incompatible orders on different threads.
- ▶ **Simplicity vs. scalability (convoying, priority inversion)?** Coarse-grained locking is simpler to program, but creates bottlenecks to kill scalability.
- ▶ **Not composable.** In today's world, this is a deadly sin.



Important assumptions

(1) **You have already measured** performance/scalability and proven you have a high-contention data structure, before resorting to the techniques described in this talk.

(2) **You will measure again** after you write a hopefully-more-concurrent replacement using these techniques to ensure that it is actually an improvement.



Roadmap

- ▶ **Two Basic Tools**
 - ▶ **Transactional thinking + `atomic<T>`**
- ▶ Basic Example: Double-Checked Locking
 - ▶ It's pretty easy to do right, but you still have to do it right
- ▶ Producer-Consumer Variations
 - ▶ Using locks, locks + lock-free, and fully lock-free
- ▶ A Singly Linked List: This Stuff Is Harder Than It Looks
 - ▶ Just find, push_front, and pop: How hard could it be?



enter critical region

exit critical region

Lock-Free Fundamental #1

► Your key concept: Think in transactions (ACID).

- Atomicity:
 - A transaction is all-or-nothing; “commit” is atomic. Other code must not be able to see the data in a partially-updated state (i.e., a corrupt state).
 - [LF] Publish each change using one atomic write (read-modify-write).
- Consistency, Isolation, Durability:
 - A transaction takes the data from one consistent state to another.
 - Two transactions never simultaneously operate on the same data.
 - A committed transaction is never overwritten by second transaction that did not see the results of the first transaction. (The “lost update” problem.)
 - [LF] Make sure concurrent updates don’t interfere with each other (especially think about deletes!) or with concurrent readers.



Lock-Free Fundamental #2

► Your key tool: The ordered atomic variable.

- C++11 “atomic<T>” and C11 “atomic_*”.
- Java “volatile T” and Atomic* (e.g., AtomicLong).
- .NET “volatile T”.
- Semantics and operations:
 - Each individual read and write is atomic, no locking required.
 - Reads/writes are guaranteed not to be reordered.
 - Compare-and-swap (CAS)... *conceptually* an atomic execution of:


```
bool atomic<T>::compare_exchange_strong( T& expected, T desired ) {
    if( this->value == expected ) { this->value = desired;  return true; }
    else /* it's not */           { expected = this->value; return false; }
}
```
 - + **compare_exchange_weak** for use in loops (is allowed to fail spuriously)
 - + **exchange** for when a “blind write” that returns the old value is sufficient
- Notes:
 - Limited to certain types that can be manipulated atomically.
 - An ‘atomic T’ may not have the same layout (e.g., alignment) as a plain T.



atomic<T> Notes

- ▶ **Lock-free** vs. lock-based implementations:
 - ▶ If T is a small type, including most built-ins, atomic<T> is implemented without locks (typically, platform-specific instructions).
 - ▶ For larger types, atomic<T> is implemented using a lock.
- ▶ **Initialization:** Remember to explicitly initialize – `atomic<int> ai{ 0 };`
- ▶ **Interleaving:** The state of the atomic<T> can change at any time between successive calls on this thread due to interleaved calls on other threads.
- ▶ **Granularity:** Logical transactions often operate on multiple objects, or on multiple calls to the same object. Example:


```
atomic<int> account1_balance = ..., account2_balance = ...;
account1_balance += amount;
account2_balance -= amount;
```

 - ▶ Those two lines still need to be externally locked, if some invariant doesn't hold in between the two calls.



Aside: Three Levels of “Lock-Freedom”

- ▶ **Wait-free (strongest, “no one ever waits”):** Every operation will complete in a bounded #steps no matter what else is going on.
 - ▶ Guaranteed system-wide throughput + starvation-freedom.
 - ▶ **Lock-free (“someone makes progress”):** Every step taken achieves global progress (for some sensible definition of progress).
 - ▶ Guaranteed system-wide throughput.
 - ▶ All wait-free algorithms are lock-free, but not vice versa.
 - ▶ **Obstruction-free (weakest, “progress if no interference”):** At any point, a single thread executed in isolation (i.e., with all obstructing threads suspended) for a bounded number of steps will complete its operation.
 - ▶ No thread can be blocked by delays or failures of other threads.
 - ▶ Doesn't guarantee progress while two or more threads run concurrently (e.g., *deadlock* is impossible, but *livelock* could be possible).
 - ▶ All lock-free algorithms are obstruction-free, but not vice versa.
- ▶ Informally, “lock-free” ≈ “doesn't use mutexes” == any of these.

Roadmap

- ▶ Two Basic Tools
 - ▶ Transactional thinking + `atomic<T>`
- ▶ **Basic Example: Double-Checked Locking**
 - ▶ **It's pretty easy to do right, but you still have to do it right**
- ▶ Producer-Consumer Variations
 - ▶ Using locks, locks + lock-free, and fully lock-free
- ▶ A Singly Linked List: This Stuff Is Harder Than It Looks
 - ▶ Just find, push_front, and pop: How hard could it be?



Bad DCL

- ▶ The pattern relies on the flag value being atomic and ordered. But an ordinary variable doesn't guarantee these properties.

```
Widget* Widget::pInstance = nullptr;
Widget* Widget::Instance() {
    if( pInstance == nullptr ) {           // 1: first check
        lock_guard<mutex> lock( mutW );    // 2: acquire lock (crit sec enter)
        if( pInstance == nullptr ) {      // 3: second check
            pInstance = new Widget();      // 4: create and assign
        }
        // 5: release lock
    }
    return pInstance;                      // 6: return pointer
}
```

- ▶ Issues:
 - ▶ **Race on pInstance: Anything can happen.** There is a race between lines 1 and 4. This could cause a wild pointer (e.g., pointer value tearing), possibly leading to random code execution.
 - ▶ **Reordering: Seeing partly-constructed objects.** If the parts of line 4 are reordered, pInstance could point to a partly-constructed object.



Correct Double-Checked Locking

- ▶ The Double-Checked Locking (DCL) pattern (un-“broken”).

```
atomic<Widget*> Widget::pInstance{ nullptr };
Widget* Widget::Instance() {
    if( pInstance == nullptr ) {           // 1: first check
        lock_guard<mutex> lock{ mutW };    // 2: acquire lock (crit sec enter)
        if( pInstance == nullptr ) {       // 3: second check
            pInstance = new Widget();       // 4: create and assign
        }
    }                                       // 5: release lock (crit sec exit)
    return pInstance;                      // 6: return pointer
}
```

- ▶ Four key points, involving both atomicity and ordering:
 - ▶ 1: Test pInstance **atomically**.
 - ▶ 2: **Then**, if that fails, take the lock.
 - ▶ 3-4a: **Then** repeat the test and construct the object.
 - ▶ 4b: **Then** assign its this pointer **atomically** to pInstance.



Slight Optimization

- ▶ This may be slightly faster (1 vs. 2 atomic loads in the main case):

```
atomic<Widget*> Widget::pInstance{ nullptr };
Widget* Widget::Instance() {
    Widget* p = pInstance;
    if( p == nullptr ) {                   // 1: first check
        lock_guard<mutex> lock{ mutW };    // 2: acquire lock (crit sec enter)
        if( (p = pInstance) == nullptr ) { // 3: second check
            pInstance = p = new Widget(); // 4: create and assign
        }
    }                                       // 5: release lock (crit sec exit)
    return p;                              // 6: return pointer
}
```

- ▶ The compiler is allowed to do this optimization for you, but:
 - ▶ it isn't required to, and
 - ▶ it's not common yet AFAIK.



Even Better: There's a Tool For That

- ▶ The general-purpose way to spell lazy initialization in C++11 is:

```
static unique_ptr<widget> widget::instance;  
static std::once_flag widget::create;  
  
widget& widget::get_instance() {  
    std::call_once( create, [=]{ instance = make_unique<widget>(); } );  
    return instance;  
}
```

- ▶ No raw *, automatic cleanup, and much lower boilerplate-to-**real code** ratio.



Best of All: There's a Tool For That

- ▶ The special-purpose way that you should use when you can (aka the Meyers Singleton!) is this:

```
widget& widget::get_instance() {  
    static widget instance;  
    return instance;  
}
```



Roadmap

- ▶ Two Basic Tools
 - ▶ Transactional thinking + `atomic<T>`
- ▶ Basic Example: Double-Checked Locking
 - ▶ It's pretty easy to do right, but you still have to do it right
- ▶ **Producer-Consumer Variations**
 - ▶ **Using locks, locks + lock-free, and fully lock-free**
- ▶ A Singly Linked List: This Stuff Is Harder Than It Looks
 - ▶ Just find, push_front, and pop: How hard could it be?



Locks and Atomics In Combination

- ▶ The key requirement is that access to a given shared mutable object is synchronized consistently...
 - ▶ Using traditional locking. (Preferred, but sometimes problematic because locks don't compose well.)
 - ▶ Using a lock-free `atomic<>` discipline. (Less deadlock, but this style tends to be really hard today.)
- ▶ ... at every given point in time.
 - ▶ It doesn't have to be the same for the lifetime of the object.
 - ▶ For example, consider handoff situations:
 - ▶ Threads 1..N share object x, synchronizing via mutex m1.
 - ▶ Then x is handed off and never looked at again by those threads.
 - ▶ Then Threads N+1..M shared x, synchronizing via mutex m2 (or via a lock-free discipline, or some other way).



Create and Publish Queue Items: 1 Producer, Many Consumers, Using Locks

► Thread 1 (producer):

```

while( ThereAreMoreTasks() ) {
    task = AllocateAndBuildNewTask();
    {
        lock_guard<mutex> lock{mut}; // enter critical section
        queue.push( task );
        cv.notify();                //
    }                               // exit critical section
}

{
    lock_guard<mutex> lock{mut}; // enter critical section
    queue.push( done );         // add sentinel; that's all folks
    cv.notify();                //
}                               // exit critical section

```



Create and Publish Queue Items: 1 Producer, Many Consumers, Using Locks

► Threads 2..N (consumers):

```

myTask = null;
while( myTask != done ) {
    {
        lock_guard<mutex> lock{mut}; // enter critical section
        while( queue.empty() )      // if not ready, don't busy-wait,
        {
            cv.wait( mut );          // release and re-enter crit sec
        }
        myTask = queue.first();      // take task
        if( myTask != done )         // remove it if not the sentinel,
            queue.pop();             // which others need to see
    }                               // exit critical section
    if( myTask != done )
        DoWork( myTask );
}

```



Quick Quiz: Where Must Those Pesky Lock-Protected Invariants Hold, Again?

- Threads 2..N (consumers):

```

myTask = null;
while( myTask != done ) {
    {
        INVARIANTS HOLD
        lock_guard<mutex> lock{mut}; // enter critical section
        while( queue.empty() )      // if not ready, don't busy-wait.
            cv.wait( mut );         // release and re-enter crit sec
        INVARIANTS HOLD
        myTask = queue.first();     // take task
        if( myTask != done )        // remove it if not the sentinel,
            queue.pop();           // which others need to see
    }
    INVARIANTS HOLD
    if( myTask != done )
        DoWork( myTask );
}

```



Questions & Answers

- Why was `mut.unlock()` not enough to exit the critical section?
 - Unlock often is enough to exit a critical section, but we have extra semantics: “We knew” that consumers are waiting on the condition variable too.
 - If we don’t `cv.notify()`, the consumers will never wake up.
- But why `cv.notify()` on only the Producer critical section exits?
 - Because the condition variable is only to notify of new additions to the queue.
 - We don’t need to wake up other consumers when we’ve taken a task away. They’re only waiting for tasks to arrive.
- Could we make unlock-and-notify a single operation by default?
 - What an interesting suggestion! Exercise for the reader...



Create and Publish Queue Items: 1 Producer, Many Consumers (Locks + LF)

- This variant uses an **atomic<Task*> head** that points to a lock-free slist, using lock-free coordination for step 1 (producer → consumers), then a lock among consumers (**sketch**):

- Thread 1 (producer):

... *build task list* ...

head = head of task queue; // publish that complete list exists

Step 1 release

- Threads 2..N (consumers) spin until the list is there, then swarm:
while(myTask == null) {

lock_guard<mutex> lock{mut};

Step 2 critical region

if(**head** != null) { // check if list exists yet

Step 1 acquire

myTask = head; // take task
head = head->next; // remove it

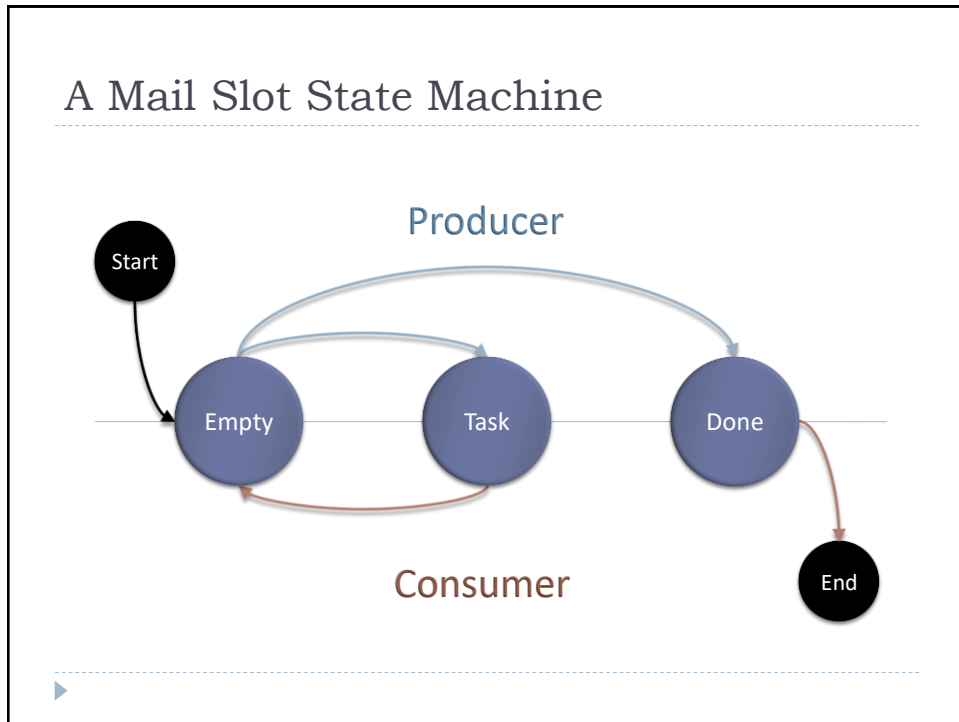
}

... = myTask->data;

- Note: In a real implementation you'd want to avoid busy-waiting.

Going Fully “Lock-Free”: Atomic Mail Slots





Create and Publish Queue Items: 1 Producer, Many Consumers, Lock-Free



- 1 Producer thread: Changes any box from null to non-null.

```

curr = 0; // keep a finger on the current mailbox
// Phase 1: Build and distribute tasks
while( ThereAreMoreTasks() ) {
    task = AllocateAndBuildNewTask();
    while( slot[curr] != null ) // acquire null: look for next empty slot
        curr = (curr+1)%K;
    slot[curr] = task; // release non-null: "You have mail!"
    sem[curr].signal();
}
// Phase 2: Stuff the mailboxes with "done" signals
numNotified = 0;
while( numNotified < K ) {
    while( slot[curr] != null ) // acquire null: look for next notifiable slot
        curr = (curr+1)%K;
    slot[curr] = done; // release done: write sentinel
    sem[curr].signal();
    ++numNotified;
}
  
```

Create and Publish Queue Items: 1 Producer, Many Consumers, Lock-Free



- ▶ K Consumer threads (mySlot = 0..K-1):
Each changes its own box from non-null to null.

```
myTask = null;
while( myTask != done ) {
    while( (myTask = slot[mySlot]) == null ) // acquire non-null,
        sem[mySlot].wait();                // without busy-wait
    if( myTask != done ) {
        slot[mySlot] = null; // release null: tell that we took it
        DoWork( myTask );    // good practice: prefer to do work
                              // outside the critical section
    }
}
```



Create and Publish Queue Items: 1 Producer, Many Consumers, Lock-Free



- ▶ K Consumer threads (mySlot = 0..K-1):
Each changes its own box from non-null to null.

```
myTask = null;
while( myTask != done ) {
    while( (myTask = slot[mySlot]) == null ) // acquire non-null,
        sem[mySlot].wait();                // without busy-wait
    if( myTask != done ) {
        slot[mySlot] = null;
        DoWork( myTask );
    }
}
```

Q: Could it make sense to swap
these two lines? Why?

// outside the critical section






Roadmap

- ▶ Two Basic Tools
 - ▶ Transactional thinking + `atomic<T>`
- ▶ Basic Example: Double-Checked Locking
 - ▶ It's pretty easy to do right, but you still have to do it right
- ▶ Producer-Consumer Variations
 - ▶ Using locks, locks + lock-free, and fully lock-free
- ▶ **A Singly Linked List: This Stuff Is Harder Than It Looks**
 - ▶ **Just find, push_front, and pop: How hard could it be?**



Example: Singly-Linked List

- ▶ A singly-linked list (aka “slist<T>”) is one of the simplest possible data structures:



- ▶ Simplifying assumptions:
 - ▶ Only four operations: Construct, destroy, find, push_front.
- ▶ **Challenge:** Write a lock-free implementation that callers can safely use without any external locks.
 - ▶ C’mon, how hard could it be?



A Lock-Free Singly-Linked List: First Cut

- ▶ Here is the interface declaration, and the internals we’ll use:

```
template<typename T>
class slist {
public:
    slist();
    ~slist();
    T* find( const T& t ) const;    // return pointer to first equal T
    void push_front( const T& t ); // insert at the front of the list
private:
    struct Node { T t; Node* next; }; // no “atomic” needed here
    atomic<Node*> head{ nullptr }; // but “atomic” is needed here:
                                   // “head” is mutable shared data

    slist(slist&) =delete;
    void operator=(slist&) =delete;
};
```



slist<T> Constructor

- ▶ The constructor is easy:

```
template<typename T>
slist<T>::slist()
{ }                                // or just "=default"
```

- ▶ Concurrency issues:

- ▶ None.
- ▶ Note: As usual, the caller has to know he can't use an object concurrently while he's constructing it. But this isn't an "external synchronization" issue as much as it's a lifetime management issue – he can't use the slist *before* it's constructed either.



slist<T> Destructor

- ▶ The destructor has to traverse:

```
template<typename T>
slist<T>::~~slist() {
    auto first = head.load(); // good habit: access head once
    while( first ) {          // (not needed here, but good habit...)
        auto unlinked = first;
        first = first->next;
        delete unlinked;
    }
}
```

- ▶ Concurrency issues:

- ▶ None.
- ▶ Note: As usual, the caller has to know he can't use an object concurrently while he's destroying it. But this isn't an "external synchronization" issue as much as it's a lifetime management issue – he can't use the slist *after* it's destroyed either.



slist<T>::find

- ▶ Return a pointer to the first equal element, or nullptr if there isn't one:

```
template<typename T>
T* slist<T>::find( const T& t ) const {
    auto p = head.load();
    while( p && p->t != t )
        p = p->next;
    return p ? &p->t : nullptr;
}
```

- ▶ Concurrency issues:
 - ▶ None.
 - ▶ As long as the constructor and destructor aren't running, this can freely run concurrently with other **find** operations... and should be safe to run concurrently with **insert** operations.



Group Exercise

- ▶ Implement push_front to insert a node with a copy of the given value:

```
template<typename T>
void slist<T>::push_front( T t ) {
```


```
}
```

- ▶ You have 10 minutes.




A Look At push_front


► Initial state:



► Intermediate state:



► Final state:



slist<T>::push_front (Flawed)

► Insert a node with a copy of the given value:

```
template<typename T>
void slist<T>::push_front( const T& t ) {
    auto p = new Node;           // create the new node
    p->t = t;                     // set its element value
    p->next = head;               // set its place in the list
    head = p;                    // publish it at the head
}
```

► Q: What's wrong with this code?



slist<T>::push_front (Flawed)

- ▶ Insert a node with a copy of the given value:


```
template<typename T>
void slist<T>::push_front( const T& t ) {
    auto p = new Node;           // create the new node
    p->t = t;                     // set its element value
    p->next = head;              // set its place in the list
    head = p;                   // publish it at the head
}
```

- ▶ **Q: What's wrong with this code?**
- ▶ **Concurrency issues:**
 - ▶ None for any readers: The insertion of the new node is atomic. A concurrent reader will see either the old value or the new value, and in either case has a valid list to traverse.
 - ▶ **Problem for writers:** What if two threads try to insert at the same time?




A Look At the Problem


- ▶ Initial state:



- ▶ Intermediate state, insertions in progress by two threads:



- ▶ Final state: First is clobbered (and leaked), last one wins.



slist<T>::push_front

- ▶ Insert a node with a copy of the given value:

```
template<typename T>
void slist<T>::push_front( const T& t ) {
    auto p = new Node;           // create the new node
    p->t = t;                     // set its element value
    p->next = head;              // set its place in the list and
    while( !head.compare_exchange_weak(p->next, p) )
        {}                      // try to swap it in until successful
}
```

- ▶ The “CAS loop” is a common construction in lock-free code.
 - ▶ Loop until “we get to be the one” to update head from ‘expected’ to ‘desired’.
- ▶ Concurrency issues:
 - ▶ None for any readers: The insertion of the new node is atomic. A concurrent reader will see either the old value or the new value, and in either case has a valid list to traverse.
 - ▶ None for writers: The CAS loop makes concurrent writers safe (for now).


Well, that was easy...

So how about adding just one more little member function?



Revised Example: Pop Goes the List

- ▶ We'll stick with our singly-linked list, one of the simplest possible data structures:



- ▶ Simplifying assumptions:
 - ▶ Original operations: Construct, destroy, find, push_front.
 - ▶ **New operation: pop** to erase the first element from the list.
- ▶ **Same challenge:** Write a lock-free implementation that callers can safely use without any external locks.
 - ▶ C'mon, how hard could it be?

▶

Group Exercise

- ▶ Implement pop_front to erase the first node:

```
template<typename T>
void slist<T>::pop_front() {
```

```
}
```

- ▶ You have 10 minutes.

▶

A Lock-Free Singly-Linked List: First Cut

- Here is the interface declaration, and the internals we'll use:


```
template<typename T>
class slist {
public:
    slist();
    ~slist();
    T* find( const T& t ) const;    // return pointer to first equal T
    void push_front( const T& t ); // insert at the front of the list
    void pop_front();              // remove first element
private:
    struct Node { T t; Node* next; }; // no "atomic" needed here
    atomic<Node*> head{ nullptr }; // but it's needed here, because
                                    // "head" is mutable shared data

    slist(slist&) =delete;
    void operator=(slist&) =delete;
};
```




A Look At pop


- Initial state:



- Intermediate state:



- Final state:



slist<T>::pop (Flawed)

- ▶ Remove the first node:

```
template<typename T>
void slist<T>::pop_front() {
    auto p = head.load();           // remember current first node
    if( p ) head = p->next;         // set head to the second node
    delete p;                       // and clean up old first node
}
```

- ▶ Q: What's wrong with this code?



slist<T>::pop (Flawed)

- ▶ Remove the first node:

```
template<typename T>
void slist<T>::pop_front() {
    auto p = head.load();           // remember current first node
    if( p ) head = p->next;         // set head to the second node
    delete p;                       // and clean up old first node
}
```

- ▶ Q: What's wrong with this code?

- ▶ Concurrency issues:

- ▶ **Problem for readers:** What if a concurrent reader doing a find() is pointing to the first node and about to read its next pointer?
- ▶ **Problem for writers:** What if a concurrent writer is trying to insert? What if a concurrent writer is trying to erase?



slist<T>::pop (Attempt #2, Still Flawed)

- ▶ Remove the first node:

```
template<typename T>
void slist<T>::pop_front() {
    auto p = head.load();           // important: read head once
    while( p && !head.compare_exchange_weak(p, p->next) )
        {}                          // NB: relies on short-circuit eval
    delete p;                       // and clean up
}
```

- ▶ Q: What's wrong with this code?



slist<T>::pop (Attempt #2, Still Flawed)

- ▶ Remove the first node:

```
template<typename T>
void slist<T>::pop_front() {
    auto p = head.load();           // important: read head once
    while( p && !head.compare_exchange_weak(p, p->next) )
        {}                          // NB: relies on short-circuit eval
    delete p;                       // and clean up
}
```

- ▶ Q: What's wrong with this code?

- ▶ Concurrency issues:

- ▶ **Same problem for readers:** What if a concurrent reader doing a find() is pointing to the first node and about to read its next pointer?
- ▶ **Subtle problem for writers:** What if a concurrent writer is trying to insert? What if a concurrent writer is trying to erase? Let's look at the "ABA problem"...




The “ABA Problem”

► Original state:




The “ABA Problem”

► Step 1 of deletion:
`p = head;`
`__temp = p->next;`



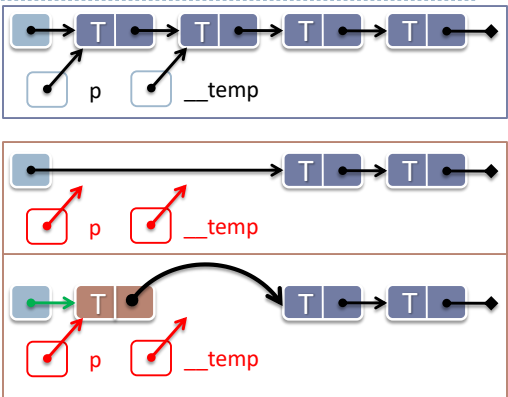
The “ABA Problem”

- ▶ Step 1 of deletion:
`p = head;`
`__temp = p->next;`
- ▶ Concurrently, another thread deletes the p two nodes.



The “ABA Problem”

- ▶ Step 1 of deletion:
`p = head;`
`__temp = p->next;`
- ▶ Concurrently, another thread deletes the p two nodes.
- ▶ Then someone inserts a new node, and the allocator **reuses the same memory block**.







The “ABA Problem”

- ▶ Step 1 of deletion:


```
p = head;
__temp = p->next;
```
- ▶ Concurrently, another thread deletes the p two nodes.
- ▶ Then someone inserts a new node, and the allocator **reuses the same memory block**.
- ▶ Step 2 (CAS) “succeeds”:


```
head.c_e_w(p, __temp);
delete p;
```

 and we’re in trouble.



ABA Solutions (sketch)

- ▶ We need to solve the ABA issue: Two nodes with the same address, but different identities (existing at different times).
- ▶ Option 1: Use lazy garbage collection.
 - ▶ Solves the problem. Memory can’t be reused while pointers to it exist.
 - ▶ But: Not an option (yet) in portable C++ code, and destruction of nodes becomes nondeterministic.
- ▶ Option 2: Use reference counting (garbage collection).
 - ▶ Solves the problem in cases without cycles. Again, avoids memory reuse.
- ▶ Option 3: Make each pointer unique by appending a serial number, and increment the serial number each time it’s set.
 - ▶ This way we can always distinguish between A and A’.
 - ▶ But: Requires an atomic compare-and-swap on a value that’s larger than the size of a pointer. Not available on all hardware & bit-nesses.
- ▶ Option 4: Use hazard pointers.
 - ▶ Maged Michael and Andrei Alexandrescu have covered this in detail.
 - ▶ But: It’s very intricate. Tread with caution.

Delete-While-Traversing Solutions (sketch)

- ▶ We also need to resolve the deletion issue: We can't delete a node if a concurrent reader/writer might be pointing to it.
 - ▶ A concurrent member function, like `find`.
 - ▶ (!) A concurrent users of a `T*` we handed out.
- ▶ Option 1: Use lazy garbage collection.
 - ▶ Solves the problem because memory can't be reused while any pointers to it exist. However, destruction of nodes becomes nondeterministic.
- ▶ Option 2: Use reference counting (garbage collection).
 - ▶ Solves the problem in cases without cycles.
- ▶ Option 3: Never actually delete a node (only logically delete).
 - ▶ Can work when deleting is rare.
- ▶ Option 4: Put auxiliary nodes in between actual nodes.
 - ▶ Contains a next pointer only, no data. These links don't move. Enables operations on adjacent nodes to run without interference.



Psst...

Interested in a correct *slist* that fell off the
back of my friend's truck, for cheap?



A Lock-Free Singly-Linked List: Second Cut

- Judicious tweaks to eliminate raw *: ref counting + *reference*.

```
template<typename T> class slist {
    struct Node { T t; shared_ptr<Node> next; };
    atomic_shared_ptr<Node> head;
    slist(slist&) =delete;
    void operator=(slist&) =delete;
```

```
public:
    slist() =default;
    ~slist() =default;
    class reference { /*...*/ };
    auto find( const T& t ) const {
        auto p = head.load();
        while( p && p->t != t )
            p = p->next;
        return reference(move(p))
    }
```

Q: How would you implement class reference?

```
class reference {
    shared_ptr<Node> p;
public:
    reference(shared_ptr<Node> p_) : p{p_} {}
    T& operator*() { return p->t; }
    T* operator->() { return &p->t; }
};
```

A Lock-Free Singly-Linked List: Second Cut

- Continued:

```
void push_front( const T& t ) {
    auto p = make_shared<Node>();
    p->t = t;
    p->next = head;
    while( !head.compare_exchange_weak(p->next, p) )
        {}
}
```

```
void pop_front() {
    auto p = head.load();
    while( p && !head.compare_exchange_weak(p, p->next) )
        {}
}

};
```

Q: Where is the "delete"?

What Happens In the Case of Concurrent...

► Pop

```
void pop_front() {
    auto p = head.load();
    while( p &&
           !head.c_e_w(p, p->next) )
        {}
}
```

► Pop

```
void pop_front() {
    auto p = head.load();
    while( p &&
           !head.c_e_w(p, p->next) )
        {}
}
```

- The only competing modifying operations are *head.compare_exchange*.
 - One will happen-before the other, and succeed.
 - The other will fail, and retry.
 - ABA can't happen because no delete+recycling.



What Happens In the Case of Concurrent...

► Insert

```
void push_front( const T& t ) {
    auto p = make_shared<Node>();
    p->t = t;
    p->next = head;
    while( !head.c_e_w(p->next, p) )
        {}
}
```

► Pop

```
void pop_front() {
    auto p = head.load();
    while( p &&
           !head.c_e_w(p, p->next) )
        {}
}
```

- The only competing modifying operations are *head.compare_exchange*.
 - One will happen-before the other, and succeed.
 - The other will fail, and retry.



What Happens In the Case of Concurrent...

Find

```
auto find( const T& t ) const {
    auto p = head.load();
    while( p && p->t != t )
        p = p->next;
    return reference(p);
}
```

Pop


```
void pop_front() {
    auto p = head.load();
    while( p &&
        !head.c_e_w(p, p->next) )
        {}
}
```

- Thanks to ref counting, *find* keeps its current node (and successors) alive.
- find* sees list “as if” *pop* waited for *find* to finish!
- Important concept: **Linearizability**.




A Look At *find* + *pop*


Initial state:



Find begins and gets partway through the list:



Two pops complete:



Translation Guide

► The code we just saw.

- From the Concurrency TS:

```
atomic_shared_ptr<T> a;
```

```
auto p = a.load();
```

```
a.compare_exchange_weak(e,d);
```

► What you actually write today.

- The “atomic<>” is a comment, and remember to write an *atomic_** call for every use of the *shared_ptr*.

```
shared_ptr<T> a;
```

```
// remember “atomic” – rely on discipline
```

```
auto p = atomic_load(&a);
```

```
atomic_compare_exchange_weak(&a,&e,d);
```

Q: How long are nodes kept alive?

A: Until the last *shared_ptr<Node>*
(including in a returned *::reference*)
to that node or an earlier *Node* goes away.

If calling code stores *::references* we can “leak” *Nodes*
in the Java sense of “leak” (delay their cleanup).

One option: Use custom refcounting to distinguish “navigating”
strong refs that keep the following *Nodes* alive, and
ordinary refs to only the *T* and that can *->next.reset()*.

A more standard option follows... ++simplicity vs. --space/time.

HT: Tomasz Kamiński

A Lock-Free Singly-Linked List: Third Cut

- + Proverbial level of indirection (NB: space/alloc overhead):

```
template<typename T> class slist {
    struct Node { shared_ptr<T> t; shared_ptr<Node> next; };
    atomic_shared_ptr<Node> head;
    slist(slist&) =delete;
    void operator=(slist&) =delete;
public:
    slist() =default;
    ~slist() =default;
    // no more 'reference' class
    auto find( const T& t ) const {
        auto p = head.load();
        while( p && p->t != t )
            p = p->next;
        return p.t;
    }
};
```



A Lock-Free Singly-Linked List: Third Cut

- Continued:

```
void push_front( const T& t ) {
    auto p = make_shared<Node>();
    p->t = make_shared<T>(t);
    p->next = head;
    while( !head.compare_exchange_weak(p->next, p) )
        {}
}

void pop_front() {
    auto p = head.load();
    while( p && !head.compare_exchange_weak(p, p->next) )
        {}
}

};
```



Roadmap

- ▶ Two Basic Tools
 - ▶ Transactional thinking + `atomic<T>`
- ▶ Basic Example: Double-Checked Locking
 - ▶ It's pretty easy to do right, but you still have to do it right
- ▶ Producer-Consumer Variations
 - ▶ Using locks, locks + lock-free, and fully lock-free
- ▶ A Singly Linked List: This Stuff Is Harder Than It Looks
 - ▶ Just find, `push_front`, and `pop`: How hard could it be?



Questions?

Bonus Slides



Example 1

- ▶ Baseline code.

```
template <typename T>
struct LowLockQueue {
    struct Node {
        Node( T val ) : value(val), next(nullptr) { }
        T value;                      // objects are held by value
        atomic<Node*> next;
    };
};
```

- ▶ *first* and *last* point to the before-the-first and last nodes
- ▶ *divider* points to a boundary between producer and consumer

```
Node *first, *last;                // for producer only
atomic<Node*> divider;              // shared: P/C boundary
atomic<bool> producerLock;         // shared by producers
atomic<bool> consumerLock;         // shared by consumers
```



Example 1 (continued)

► Construct.

```
public:
    LowLockQueue() {
        first = divider = last = new Node( T() );
        producerLock = consumerLock = false;
    }
```

► Destroy.

```
~LowLockQueue() {
    while( first != nullptr ) {
        Node* tmp = first;
        first = tmp->next;
        delete tmp;
    }
}
```

Example 1 (continued)

► *Consume* returns the value in the first unconsumed node.

- Note: The entire body of *Consume* is inside the critical section, so we get no concurrency among consumers in this code.

```
bool Consume( T& result ) {
    while( consumerLock.exchange(true) )
    {
        // acquire exclusivity
        if( divider->next != nullptr ) { // if queue is nonempty
            result = divider->next->value; // copy it back to the caller
            divider = divider->next; // publish that we took an item
            consumerLock = false; // release exclusivity
            return true; // and report success
        }
        consumerLock = false; // release exclusivity
        return false; // queue was empty
    }
}
```

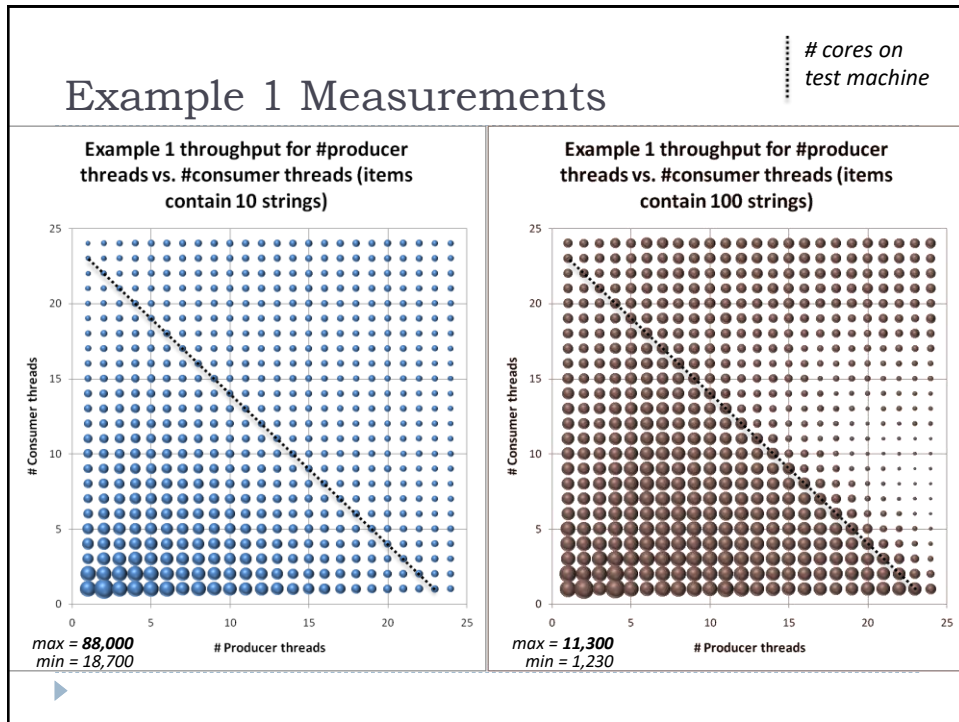
Example 1 (continued)

- ▶ *Produce* adds a new node to the tail, then lazily cleans up any consumed nodes at the front of the list.
- ▶ Note: Not all of the body of *Produce* is inside the critical section, so there is some concurrency among producers in this code.

```
bool Produce( const T& t ) {
    Node* tmp = new Node( t );    // do work off to the side
    while( producerLock.exchange(true) )    // acquire exclusivity
    {
        last = last->next = tmp;    // publish the new item
        while( first != divider ) {    // lazy cleanup
            Node* tmp = first;
            first = first->next;
            delete tmp;
        }
        producerLock = false;    // release exclusivity
        return true;
    }
};
```

How Fast Is It?

- ▶ Key properties to look for:
 - ▶ **Throughput:** Total work, here #objects that can pass through the queue.
 - ▶ **Scalability:** Ability to use more hardware (cores) to get more work done.
- ▶ These are affected by the effects of:
 - ▶ **Contention:** How much threads interfere with each other by fighting for resources.
 - ▶ **Oversubscription:** What happens if there is more CPU-bound work ready to execute than available hardware to execute it.



Ex. 2: Shrinking Consumer Critical Section

- ▶ For better performance, add heap-allocation...
... *what, what?*
 - ▶ Lets us move the copying work out of the critical section.
- ▶ Differs from Example 1 (part 1 of 2):

```
struct Node {
    Node( T* val ) : value(val), next(nullptr) { }
    T* value;
    atomic<Node*> next;
};

LowLockQueue() {
    first = divider = last = new Node( nullptr );
    producerLock = consumerLock = false;
}
```

Ex. 2: Shrinking Cons Crit Sec (continued)

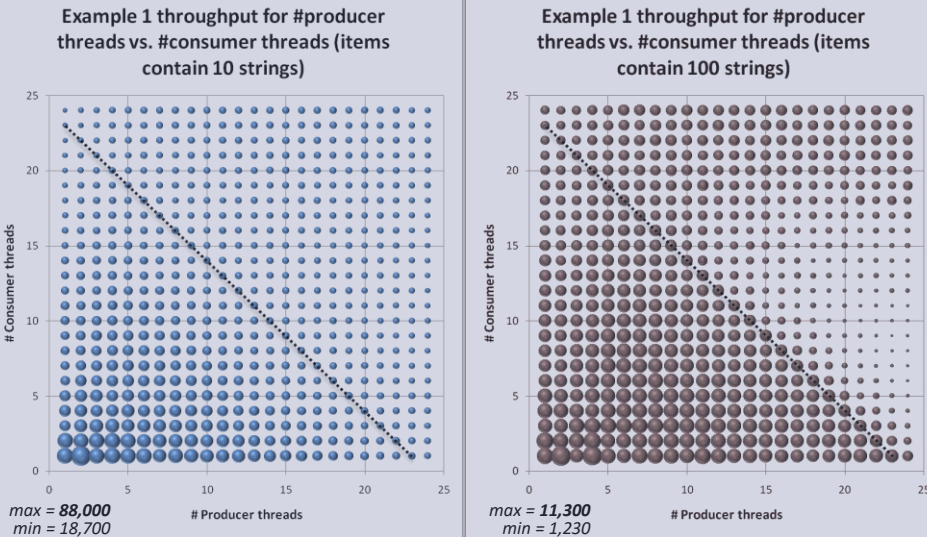
- Differs from Example 1 (part 2 of 2):
 - Move the copying of the dequeued object, and the deletion of the value, outside the critical section.

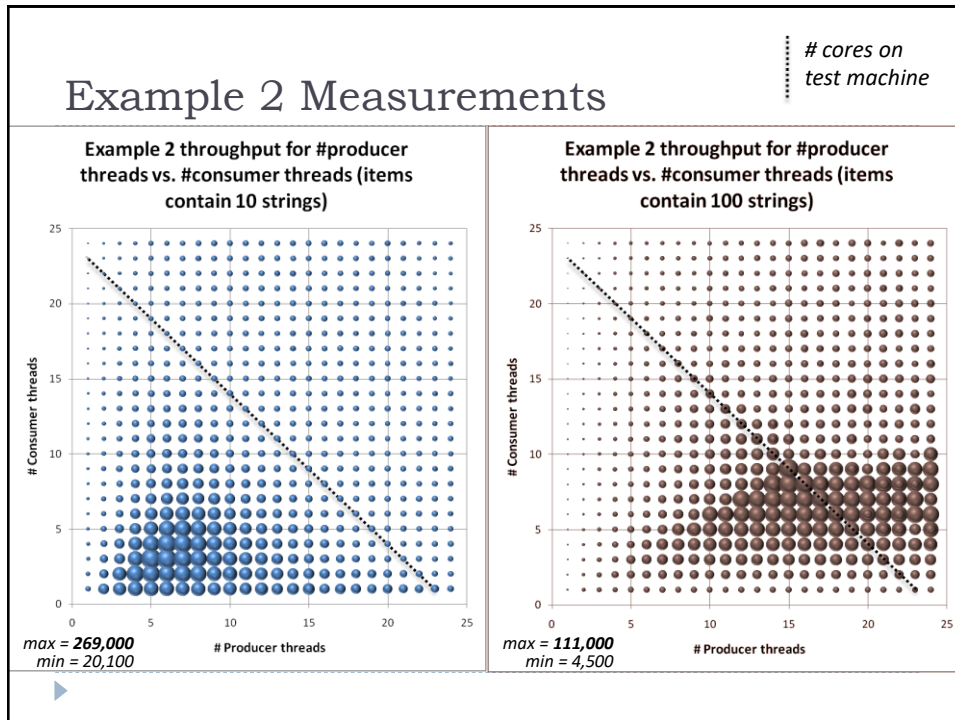
```
bool Consume( T& result ) {  
    while( consumerLock.exchange(true) )  
    {  
        // acquire exclusivity  
        if( divider->next != nullptr ) {  
            // if queue is nonempty  
            T* value = divider->next->value;  
            // take it out  
            divider->next->value = nullptr;  
            // of the Node  
            divider = divider->next;  
            // publish that we took an item  
            consumerLock = false;  
            // release exclusivity  
            result = *value;  
            // now copy it back to the caller  
            delete value;  
            return true;  
            // and report success  
        }  
        consumerLock = false;  
        // release exclusivity  
        return false;  
        // queue was empty  
    }  
}
```



Example 1 Measurements

cores on
test machine





Ex. 3: Reducing Head Contention

- ▶ Ex. 1 & 2: Producer lazily removed consumed nodes.
 - ▶ Forces producer to touch both ends of the queue.
 - ▶ All threads (producers and consumers) have to touch head.
 - ▶ Even though producers and consumers use different locks and can run concurrently w.r.t. each other, this results in invisible contention in the memory system.
- ▶ Idea: Let each consumer trim the nodes it consumed.
 - ▶ Which it was touching anyway \Rightarrow better **locality**.
 - ▶ Bonus: No more *divider*.
- ▶ Differs from Example 3 (part 1 of 3):

```
LowLockQueue() {
    first = last = new Node( nullptr ); // no more divider
    producerLock = consumerLock = false;
}
```

Ex. 3: Reducing Head Contention (cont'd)

- Differs from Example 3 (part 2 of 3):

```
bool Consume( T& result ) {
    while( consumerLock.exchange(true) )
        {} // acquire exclusivity
    if( first->next != nullptr ) { // if queue is nonempty
        Node* oldFirst = first;
        first = first->next;
        T* value = first->value; // take it out
        first->value = nullptr; // of the Node
        consumerLock = false; // release exclusivity
        result = *value; // now copy it back
        delete value; // and clean up
        delete oldFirst; // both allocations
        return true; // and report success
    }
    consumerLock = false; // release exclusivity
    return false; // queue was empty
}
```



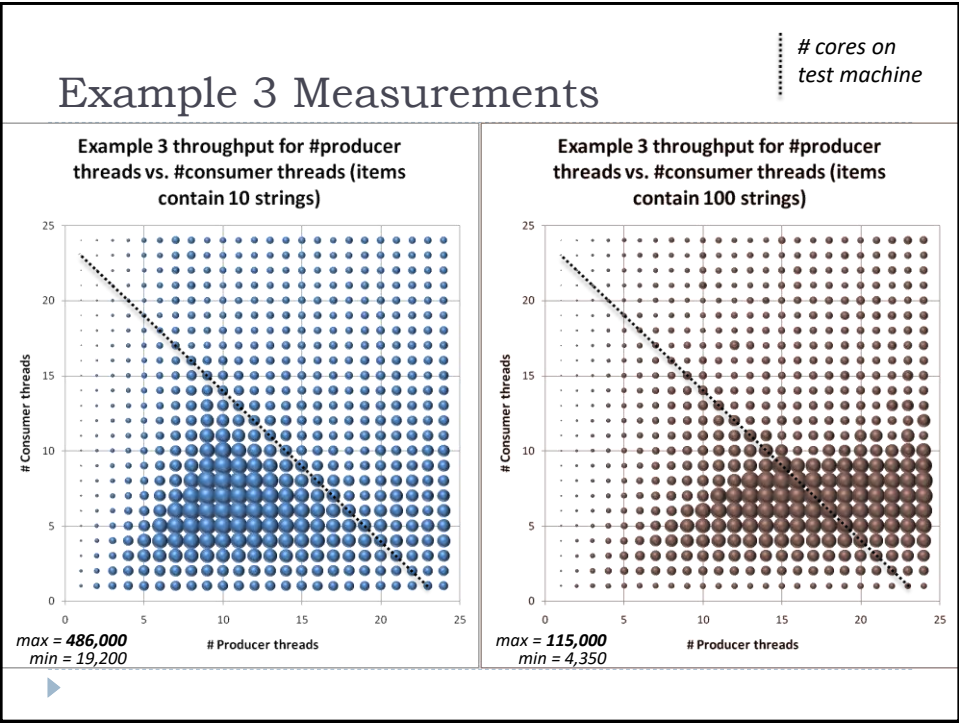

Ex. 3: Reducing Head Contention (cont'd)

- Differs from Example 3 (part 3 of 3):

- *Producer* is simpler.

```
bool Produce( const T& t ) {
    Node* tmp = new Node( t ); // do work off to the side
    while( producerLock.exchange(true) )
        {} // acquire exclusivity
    last->next = tmp; // A: publish the new item
    last = tmp; // B: not "last->next"
    producerLock = false; // release exclusivity
    return true;
}
```





Ex. 4: Do Nothing... or, “Add Nothing”

- ▶ Keep data that is *not* used together *apart*.
 - ▶ If variables A and B are liable to be used on different threads, keep them on separate cache line.
- ▶ Differs from Example 4:

```

struct Node {
    Node( T* val ) : value(val), next(nullptr) { }
    T* value;
    atomic<Node*> next;
};

Node* first;
atomic<bool> consumerLock;
Node* last;
atomic<bool> producerLock;

```



Ex. 4: Do Nothing... or, “Add Nothing”

- ▶ Keep data that is *not* used together *apart*.
 - ▶ If variables A and B are liable to be used on different threads, keep them on separate cache line.
- ▶ Differs from Example 4:

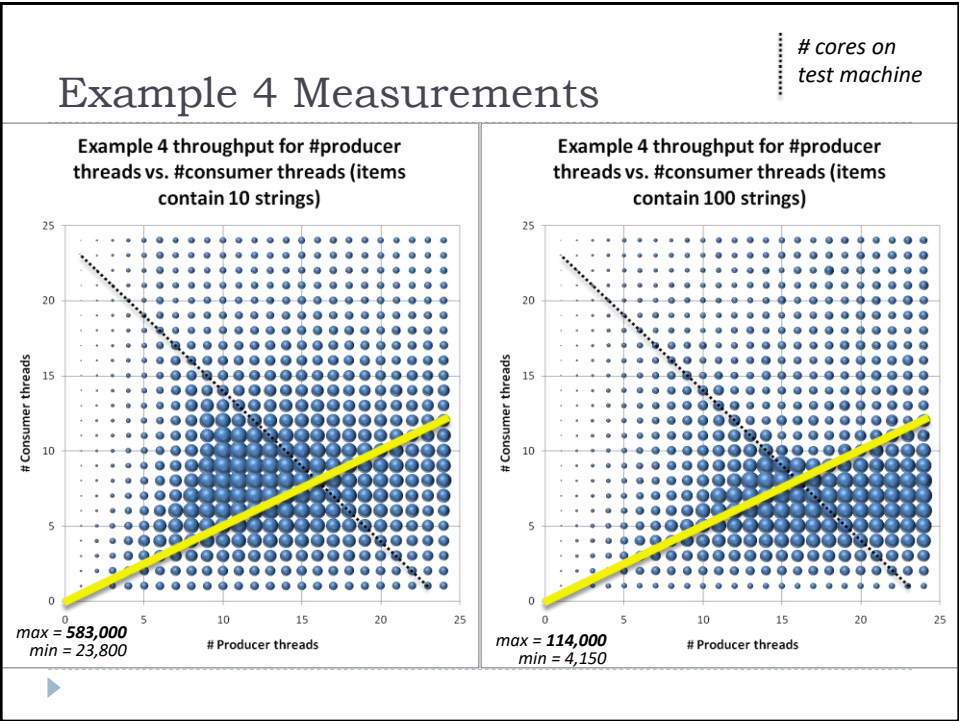

```

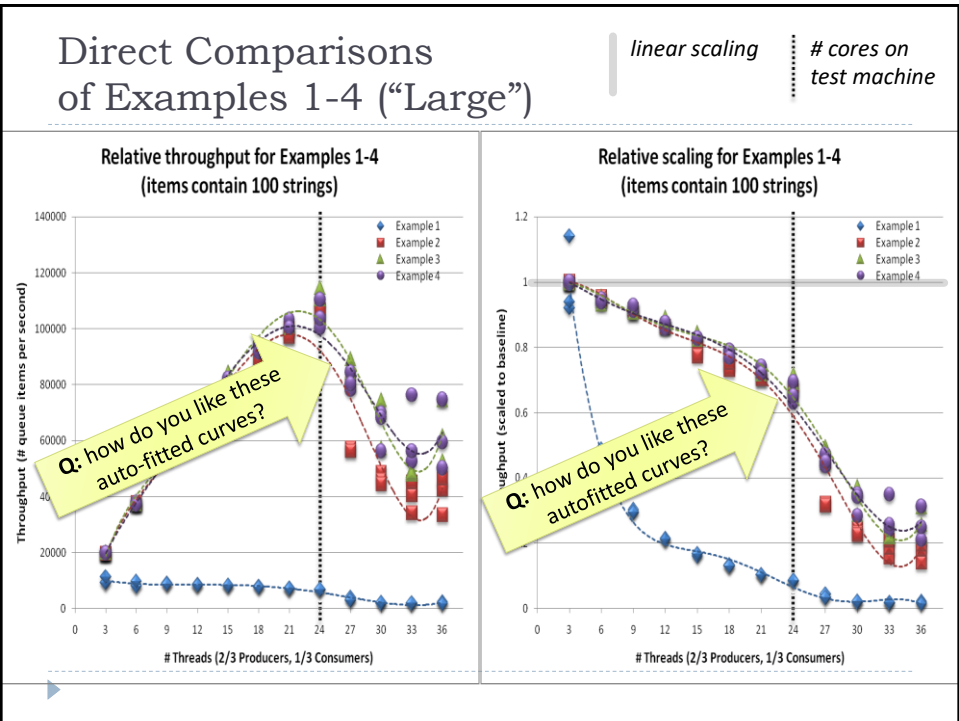
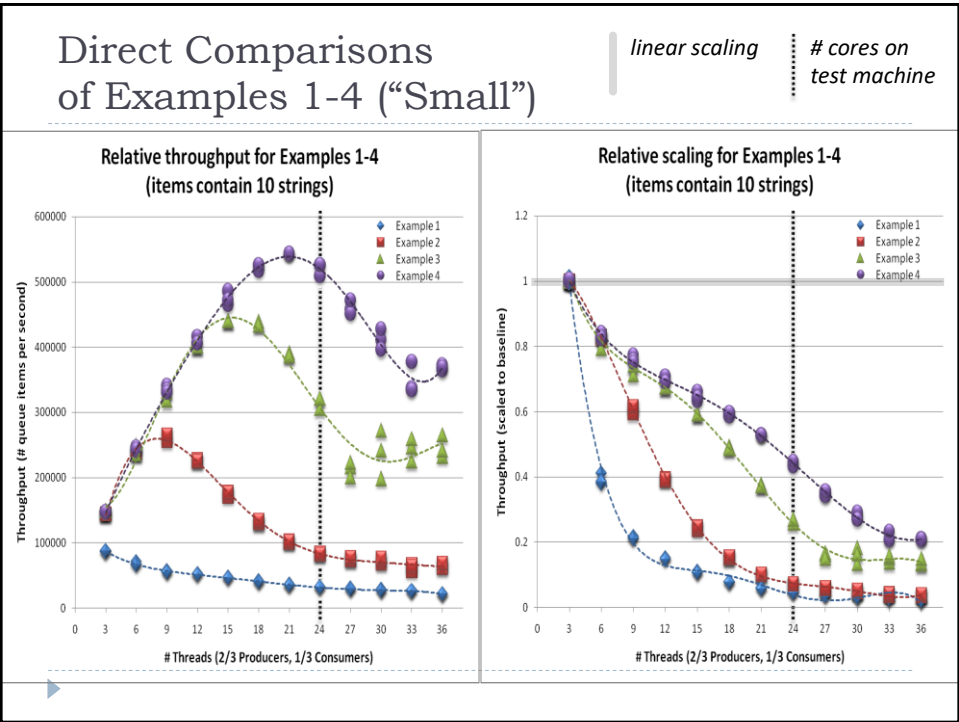
struct alignas(CACHE_LINE_SIZE) Node {
    Node( T* val ) : value(val), next(nullptr) { }
    T* value;
    atomic<Node*> next;
};

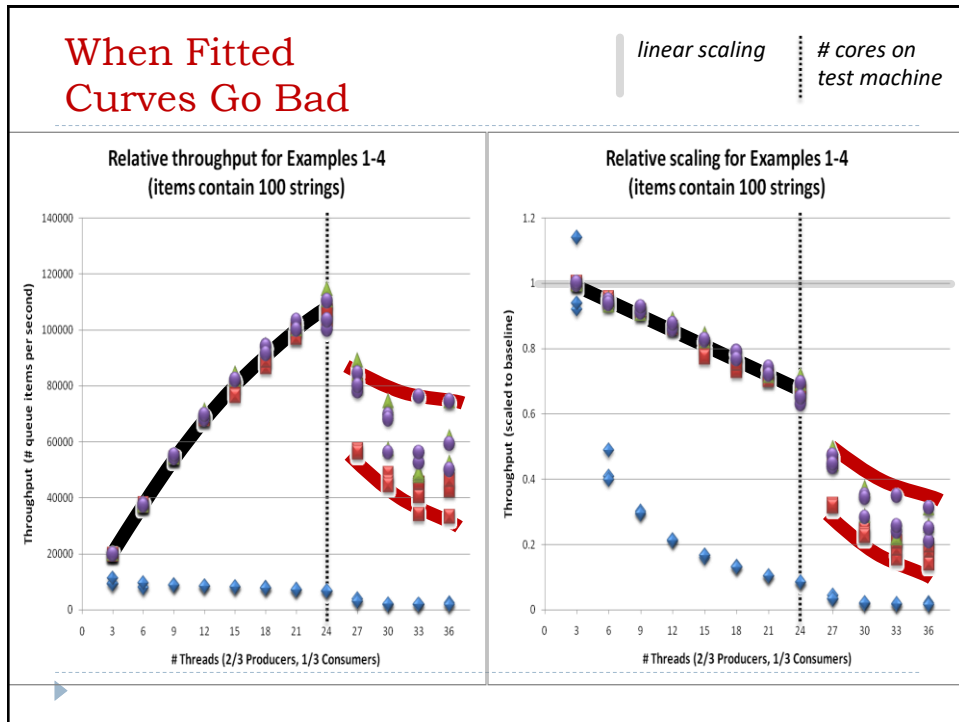
alignas(CACHE_LINE_SIZE) Node* first;
alignas(CACHE_LINE_SIZE) atomic<bool> consumerLock;
alignas(CACHE_LINE_SIZE) Node* last;
alignas(CACHE_LINE_SIZE) atomic<bool> producerLock;

```









What Have We Learned?

- ▶ To improve **scalability**, we need to minimize **contention**:
 - ▶ Reduce the size of critical sections \Rightarrow more concurrency.
 - ▶ Reduce sharing by isolating threads to use different parts of the data structure.
 - ▶ Moving cleanup from producer to consumer lets consumers touch only the head, producers touch only the tail.
 - ▶ Reduce false sharing of different data on the same cache line, but adding alignment padding.
 - ▶ Separate variables that should be able to be used concurrently by different threads should be far enough apart in memory.

What Have We Learned? (2)

- ▶ To understand scalability, need to know what to measure:
 - ▶ **Identify the key different kinds of work:** Here, producer threads and consumer threads. **Use stress tests to measure** the impact of having different quantities and combinations of these in our workload.
 - ▶ **Identify the different kinds of data:** Here, representative “small” and “large” queue items). **Vary those** to measure their impact.
 - ▶ **Measure total throughput**, or items handled per unit time.
 - ▶ **Look for scalability**, or the change in throughput as we add more threads. Does using more threads do more total work? Why or why not? In what directions, and for what combinations of workloads?
 - ▶ **Look for contention**, or the interference between multiple threads trying to do work concurrently.
 - ▶ **Watch for the cost of oversubscription**, and eliminate it either algorithmically or by limiting the actual amount of concurrency to avoid it altogether.
 - ▶ **Beware of overreliance on automated trendlines / fitted curves.** Apply them only after first examining the raw data.

Questions?